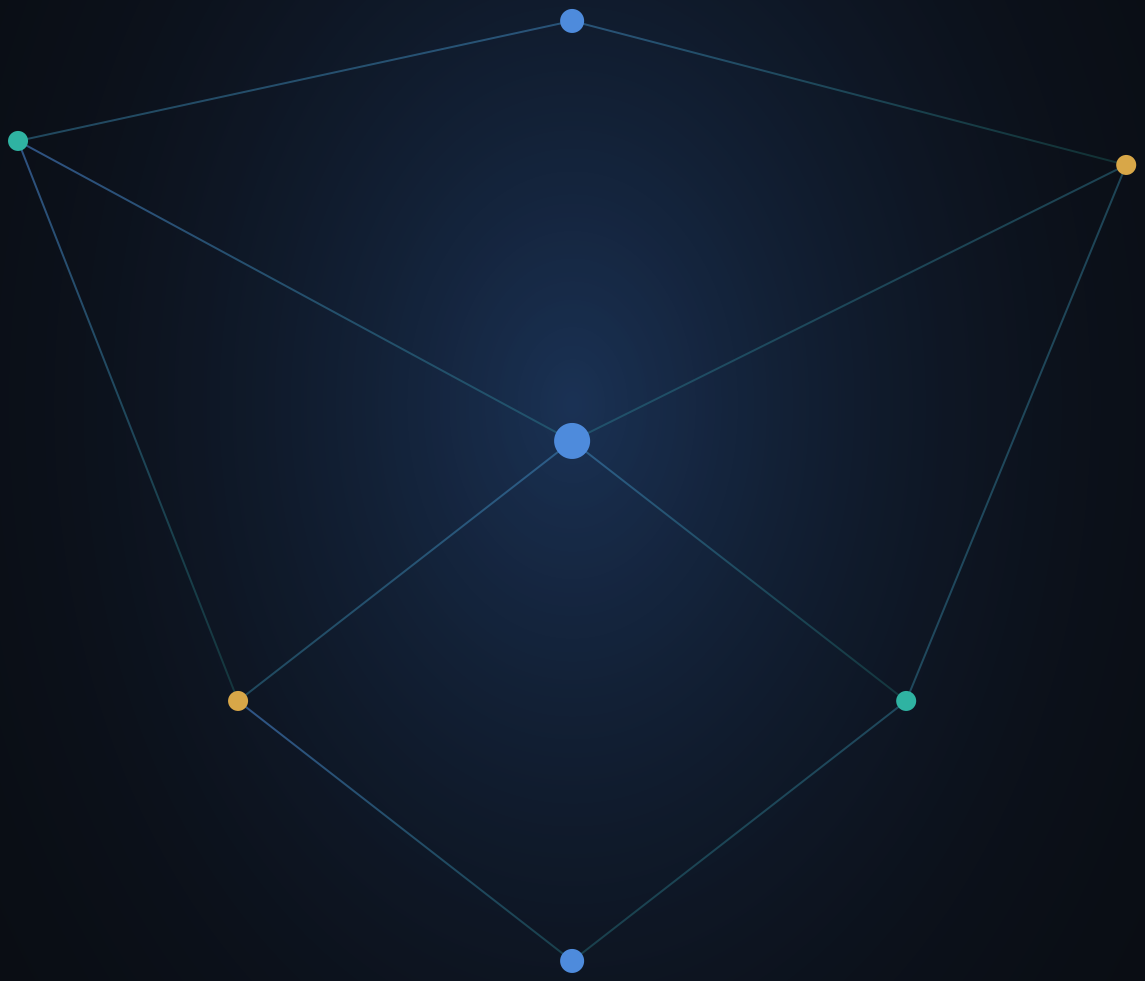
Data Intelligence OS



Confidential · Product Vision Document

Data Intelligence OS

The intelligence layer for enterprise data

An AI-powered platform that helps organisations discover, understand, trust, govern and prepare data for analytics and AI — without replacing the systems they already run.

VERSION	DATE	PREPARED BY	CLASSIFICATION
1.0	October 2026	PyrunAI Services LLP	Confidential

CONTENTS

01 Executive Summary	02	**09** Product Workflow	11
02 Product Vision	03	**10** Key Platform Modules	12
03 Problem Statement	04	**11** AI Agent Strategy	14
04 Market Opportunity	05	**12** Competitive Landscape	15
05 Target Users	06	**13** Unique Value Proposition	16
06 Target Industries	08	**14** Success Metrics	17
07 Product Positioning	09	**15** Roadmap Vision	18
08 Core Product Pillars	10	**16** Final Vision Statement	19

SECTION 01

Executive Summary

Data Intelligence OS is an AI-powered platform that sits above an organisation's existing data systems and answers one question they cannot answer today: **can this dataset be trusted?**

What it is

A layer that connects to existing warehouses and databases, discovers what is there, measures its quality, finds and verifies relationships, assembles datasets with the arithmetic proven correct, documents them automatically and tracks every change. It runs on top of SQL Server, PostgreSQL, Snowflake, Databricks and others — it replaces none of them.

Why it exists

Enterprises have spent a decade buying storage and compute. They can hold any volume of data and query it in seconds. What they still cannot do is tell a board member, with evidence, that the number on the slide is right. Every organisation has a story about a figure that was wrong for months before anyone noticed.

The business problem

Building a trusted dataset is slow, manual and fragile. A request that should take two days takes six weeks: finding the right tables, guessing how they join, cleaning the values, discovering mid-way that a join multiplied the amounts, reconciling the totals by hand, and documenting none of it. When the engineer who built it leaves, the reasoning leaves with them. The dataset survives; the understanding does not.

TARGET **6×** faster dataset delivery

TARGET **100%** of published datasets reconciled

TARGET **Zero** undocumented published datasets

TARGET **Full** decision history retained

Why this is not ETL, and not AutoML

An ETL tool moves data and assumes the instructions it was given are correct. An AutoML tool trains models and assumes the dataset it was handed is correct. Data Intelligence OS exists in the gap between them: it is the system that establishes whether the dataset *is* correct, explains how it reached that conclusion, and refuses to publish when it is not. That refusal — a reconciliation gate no reviewer can override — is the product's defining behaviour.

SECTION 02

Product Vision

> Every dataset an enterprise uses should arrive explained, reconciled and signed off — and the reasoning behind it should outlive the person who built it.
>
> Vision statement · Data Intelligence OS

Trusted data

A dataset is not trusted because a tool produced it. It is trusted because the totals were compared against the source and matched, and because the comparison is on record. Trust is evidence, not reputation.

AI-assisted engineering

AI narrows the search space: twelve candidate joins out of two thousand column pairs, a first-draft query, a drafted description. A person decides. The platform never moves from assistant to author.

Data understanding

Schema tells you a column is a DECIMAL. Understanding tells you it is a sanctioned amount in rupees, excluding written-off exposure, as defined by Credit Risk in version 7.

Dataset creation

Assembling a dataset should be a reviewed, versioned, repeatable act — not a one-off query someone pastes into a notebook and forgets.

Governance

Governance that is a by-product of doing the work, not a separate system someone fills in afterwards and nobody maintains.

Data intelligence

The platform accumulates knowledge: why a join exists, why a rule was added, why a suggestion was rejected. That record is the asset.

The long-term position: an intelligence layer, not another store

The industry has consolidated storage. Snowflake, Databricks and the cloud warehouses have won that layer and will keep it. Competing there is futile. The layer that remains unclaimed sits above them — the one that knows what the data *means*, whether it is *right*, and *who decided* so.

LAYER	OWNED BY	ANSWERS THE QUESTION	OUR POSITION
Storage & compute	Snowflake, Databricks, cloud warehouses	Where does the data live and how fast can I query it?	Build on it
Movement	Fivetran, Informatica, dbt, in-house ETL	How does data get from A to B?	Integrate with it
Cataloguing	Atlan, Collibra, Alation, DataHub	What exists and who owns it?	Interoperate, not replace
Intelligence & trust	Largely unclaimed	Is this dataset correct, and how do we know?	**Own it**

What we deliberately exclude. Model training is out of scope. It is a different product with different buyers and a crowded field. Our responsibility ends the moment a dataset is proven ready — and that boundary is what keeps the product coherent.

SECTION 03

Problem Statement

The problems below are not hypothetical. Each one is a pattern the panel has seen repeatedly inside mid-to-large enterprises with mature data platforms.

Data scattered across systems

A single customer view needs core banking, a CRM, a payments warehouse, a complaints database and three spreadsheets. No one person has access to all five.

Thousands of undocumented tables

A typical warehouse holds 2,000–10,000 tables. Fewer than 15% carry a usable description. Table names look like `T_CUST_M_01` and `ZFI_BSEG_X`.

Unknown relationships

Foreign keys were dropped years ago for load performance. Two columns called `customer_id` often hold two different identifiers from two different eras.

Slow dataset creation

A request that should take two days routinely takes four to six weeks, most of it spent discovering what exists and verifying that a join behaves.

Business depends on data engineers

Every question becomes a ticket. The engineering queue becomes the organisation's rate limit on decision-making.

Lack of trust in numbers

Three reports show three different revenue figures. Each is defensible. The meeting is spent reconciling rather than deciding.

The silent failure that defines the category. A loan table joined to its repayment history multiplies each loan by the number of EMI rows. Summing the sanctioned amount after that join returns a figure roughly thirty-nine times too large. Nothing errors. The query succeeds. The report renders. In the cases we have examined, the error survived for weeks because no system compared the output total against the source total.

Knowledge loss

The engineer who knew that pre-2019 records follow different rules resigns. The rule was never written down, so the next person rebuilds the dataset and reintroduces the bug.

Lineage that stops

Lineage diagrams terminate at a four-thousand-line stored procedure. The picture looks complete and is missing its middle.

Impact unknown

"If I change this column, what breaks?" is answered by asking around. The answer usually arrives after the change, as an incident.

WORKED EXAMPLE — RETAIL BANK, 4.2 M CUSTOMERS

STAGE	TODAY	WHY IT TAKES THAT LONG
Find the tables	5–8 days	No reliable catalogue; relies on asking three teams
Work out the joins	4–6 days	No declared keys; verification is manual sampling
Clean, standardise, find the fan-out error	8–25 days	1,942 spellings of 684 city names is a normal finding; the fan-out is often caught only when a business user disputes a number
Reconcile totals, then document	3–5 days	Reconciliation is manual and frequently skipped; documentation rarely happens at all

SECTION 04

Market Opportunity

Published analyst estimates for this category vary by an order of magnitude depending on how the boundary is drawn. We present the published anchors with their sources, then a bottom-up estimate we consider more defensible.

DATA GOVERNANCE **~$6 B** global market projected for 2026 (Global Industry Analysts)

ENTERPRISE METADATA MGMT **~23%** CAGR projected to 2034, reaching ~$84 B (Market.us)

OUR READ **Wide** estimates differ by 10× — treat all as directional only

An honest caveat. The spread between these published figures is large enough that neither should anchor an investment case on its own. The category boundary — catalogue, governance, observability, quality, metadata — is drawn differently by every firm. We therefore build our case bottom-up.

BOTTOM-UP ADDRESSABLE MARKET

SEGMENT	ORGANISATIONS	REALISTIC ANNUAL CONTRACT	SEGMENT VALUE
India — banks, NBFCs, insurers	~900 regulated entities with a data function	₹18–60 lakh	₹160–540 Cr
India — pharma, manufacturing, telecom	~1,400	₹12–40 lakh	₹170–560 Cr
South-East Asia & Middle East — regulated	~2,000	$30–90 K	$60–180 M
US / EU mid-market regulated	~12,000	$40–150 K	$480 M–1.8 B

Why the market is moving now

- **AI raised the cost of bad data.** A wrong figure in a report is embarrassing. The same figure in a model's training set is systematic and invisible.
- **Regulators began asking for provenance.** Increasingly the question is not what the number is, but how it was produced.
- **Catalogue disillusionment.** Many enterprises bought a catalogue, populated 20% of it and quietly stopped. The appetite is now for tools that do work, not tools that describe work.
- **Team sizes are flat while data volume is not.** Leverage must come from tooling.

Why the market is hard

- **Budgets are committed.** Governance spend is often locked into multi-year contracts.
- **Platform vendors are absorbing adjacent layers.** Anything bundled into a warehouse contract is effectively free to the buyer.
- **Security review is a gate, not a formality.** SOC 2, penetration testing and data residency are prerequisites, not differentiators.
- **Sales cycles in regulated industries run 9–18 months.** This shapes funding requirements more than product scope does.

SECTION 05

Target Users

Eight personas, in order of how directly the platform serves them. The first four hold the daily pain; the second four hold the budget and the accountability.

Primary user

Data Engineer

RESPONSIBILITIES

Builds and operates pipelines. Owns production reliability, source integration and the correctness of what reaches the warehouse.

CURRENT CHALLENGES

Spends more time discovering and verifying than building. Carries undocumented knowledge in their head. Is the escalation point for every disputed number.

HOW WE HELP

Profiling in minutes rather than days. Join candidates measured on live data. Automatic reconciliation, which is the task most often skipped. Documentation generated rather than written.

Primary user

Analytics Engineer

RESPONSIBILITIES

Turns raw tables into modelled, tested, documented datasets the business consumes. Usually owns the dbt project.

CURRENT CHALLENGES

Tests catch what was anticipated. Nothing catches the grain error nobody thought to test for, or an upstream type change that passes silently.

HOW WE HELP

Grain and fan-out checked automatically before a model ships. Contract breaks surfaced at load. Lineage derived from the compiled plan rather than maintained by hand.

Consumer

Business Analyst

RESPONSIBILITIES

Answers business questions, builds reports, defends numbers in front of stakeholders.

CURRENT CHALLENGES

Waits weeks for datasets. Cannot tell whether a table is authoritative or an abandoned copy. Discovers definitions differ between reports only when challenged.

HOW WE HELP

A searchable catalogue of certified datasets with owners and definitions attached. Confidence that a published dataset has already been reconciled.

Consumer

Data Scientist

RESPONSIBILITIES

Frames problems, engineers features, builds and validates models, explains them to risk committees.

CURRENT CHALLENGES

Loses 40–60% of project time to data preparation. Discovers leakage late. Cannot reproduce a dataset from six months ago.

HOW WE HELP

Versioned, reproducible datasets with full lineage. Leakage warnings on candidate features. Explicit readiness status before training begins.

Target Users — decision makers

Influencer — Data Architect

RESPONSIBILITIES

Designs the data landscape, sets standards, decides which systems hold which truth, approves new tooling.

CURRENT CHALLENGES

The documented architecture and the real one diverged years ago. Standards exist on paper and are unenforced in practice.

HOW WE HELP

A live map of the real landscape derived from the systems themselves, with relationships measured rather than declared, so the architecture document reflects reality.

Economic buyer — Chief Data Officer

RESPONSIBILITIES

Accountable for data strategy, quality, governance and — increasingly — for whether the organisation's AI programmes rest on sound foundations.

CURRENT CHALLENGES

Cannot evidence data quality to the board in numbers. Has often already invested in a catalogue that stalled below 30% coverage.

HOW WE HELP

Measurable quality, reconciliation and documentation coverage as reportable metrics, produced as a by-product of work rather than a separate programme.

Approver — Governance Team

RESPONSIBILITIES

Maintains definitions, ownership, policies and stewardship across domains.

CURRENT CHALLENGES

Governance is a parallel activity nobody updates. Definitions live in a wiki that no system reads.

HOW WE HELP

Definitions are stored once, versioned, owned, and read by the platform *before* the schema — so the glossary becomes operational rather than documentary.

Approver — Compliance Team

RESPONSIBILITIES

Evidences control to regulators and auditors. Answers who accessed what, who changed what, and on what authority.

CURRENT CHALLENGES

Evidence is reconstructed after the fact from logs that were never designed to answer the question.

HOW WE HELP

An immutable record of who changed what, when and **why** — with the reason captured at the moment of the action rather than reconstructed later.

Buyer reality. The Data Engineer and Analytics Engineer feel the pain most acutely but rarely hold budget. The Chief Data Officer holds budget but has often already committed it. Our commercial strategy therefore leads with a paid pilot owned by Head of Data Engineering, sized to prove value inside one quarter, with CDO sponsorship secured on evidence rather than on promise.

SECTION 06

Target Industries

Fit is strongest where a wrong number carries a regulatory or financial consequence. Industries that tolerate approximate reporting are poor targets regardless of their data volume.

INDUSTRY	PRIORITY	WHY THEY BENEFIT
Banking	Tier 1	Regulatory returns must reconcile. A restated figure is a reportable event. Customer, loan, transaction and branch data sit in four systems that must agree.
NBFC	Tier 1	Same regulatory pressure with smaller data teams, so leverage from tooling matters more. Collections and delinquency data is notoriously inconsistent.
Insurance	Tier 1	Policy, claim and premium data is effective-dated. Joining without an as-of filter silently multiplies rows — exactly the failure the platform prevents.
Healthcare	Tier 1	Patient and clinical data carries the strictest provenance and access requirements. Lineage and audit are mandatory, not optional.
Manufacturing	Tier 2	Sensor, MES and ERP data must be reconciled before yield and OEE figures mean anything. SAP table naming makes discovery genuinely hard.
Telecom	Tier 2	Very high volume, heavy regulatory reporting, and revenue assurance is an established, funded discipline that this product directly serves.
Logistics	Tier 2	Shipment, route and billing data spans many partner systems. Cross-system reconciliation is the core operational problem.
Retail	Tier 3	Large volumes, but tolerance for approximate figures is higher and regulatory consequence is lower. Pursue opportunistically, not strategically.

Where we deliberately do not compete

Technology companies build this internally and will not buy it. Early-stage and small organisations with fewer than fifty tables do not have the problem. Marketing and growth analytics teams optimise for speed over correctness, which inverts our value proposition.

The qualifying question

One question separates a prospect from a tourist: *"When a published number turns out to be wrong, what happens?"* If the answer involves a regulator, a restatement or a board, they are a buyer. If the answer is "we fix it quietly," they are not.

SECTION 07

Product Positioning

> For data and analytics teams in regulated enterprises who cannot prove that the datasets they publish are correct, Data Intelligence OS is an AI data intelligence platform that discovers, verifies, assembles and documents trusted datasets on top of the systems they already run.
>
> Unlike catalogues, which describe data, and unlike ETL tools, which move it, Data Intelligence OS proves the result is right and refuses to publish when it is not.
>
> Positioning statement

What we are

- ✓ **AI Data Intelligence Platform** — understanding and trust as the product
- ✓ **AI-Assisted Data Engineering Platform** — the engineer stays in command
- ✓ **Trusted Dataset Platform** — reconciliation as a blocking gate
- ✓ **Enterprise Data Understanding Platform** — meaning, not just metadata

What we are not

- ✗ **An AutoML tool** — model training is explicitly out of scope
- ✗ **A traditional ETL tool** — we do not compete with Informatica or Fivetran on movement
- ✗ **A BI tool** — we produce the dataset; Power BI and Tableau visualise it
- ✗ **A data warehouse** — we store no system of record

THE ONE-SENTENCE TEST

If a prospect can describe us in one sentence as *"the tool that proves our datasets are right,"* the positioning is working. If they describe us as *"another catalogue"* or *"an AI thing for data,"* it is not — and the deal will be lost on a feature comparison against a better-funded incumbent.

Strategic discipline. The temptation in this category is to widen: add cataloguing, add ingestion, add observability, add modelling. Every added module invites comparison with a specialist who does it better. Our defensibility comes from depth in verification — grain correctness, fan-out handling, reconciliation and decision provenance — not from breadth of surface area.

SECTION 08

Core Product Pillars

Eight pillars. The first three establish understanding, the middle two produce the dataset, and the last three make it durable, explainable and governed.

PILLAR 01

Data Discovery

Connect a source and map everything in it — tables, columns, types, volumes, declared keys and candidate relationships — without a human writing an inventory.

PILLAR 02

Data Understanding

Move from structure to meaning: what a column represents in business terms, which definition governs it, and which of several similar tables is authoritative.

PILLAR 03

Data Quality

Six measured dimensions — completeness, uniqueness, validity, consistency, freshness and accuracy — with rules a domain owner can set and a score that is explainable.

PILLAR 04

Dataset Engineering

Assemble datasets with the grain declared and enforced, measure tables aggregated before joining, and fan-out treated as a compile-time error rather than a surprise.

PILLAR 05

Trust & Validation

Compare every published measure against its source within tolerance. A variance outside tolerance blocks release and cannot be waived by a reviewer. This is the defining pillar.

PILLAR 06

Documentation

Generated from the catalogue and the build record, so it cannot drift from what the pipeline actually does. Documentation that is a by-product, not a chore.

PILLAR 07

Lineage & Impact

Column-level lineage derived from the compiled plan rather than inferred, read forwards for provenance and backwards to answer what a proposed change will break.

PILLAR 08

Governance

Who changed what, when and why — captured at the moment of the decision. Policy enforcement, access control and sensitive-data classification as operational controls.

The pillar that matters most

Pillar 05. Discovery, quality and documentation are table stakes that several vendors address competently. A reconciliation gate that refuses to publish an understated figure is rare, concrete, demonstrable in a two-minute demo, and directly connected to a cost the buyer has already paid.

How the pillars reinforce one another

Discovery feeds quality. Quality constrains engineering. Engineering produces the lineage that makes impact analysis possible. Documentation and governance are generated from the same build record. No pillar is a standalone module bolted on — each one consumes the output of the last.

SECTION 09

Product Workflow

Ten stages from connection to a dataset that is ready for analytics and AI. Two of the ten are gates: validation can warn, reconciliation can block.

1. **Connect data**
   Read-only connection to the source system, credentials held in the tenant vault.
2. **Discover data**
   Tables, columns, types, volumes and declared keys inventoried automatically.
3. **Understand relationships**
   Candidate joins proposed, then measured on live data. Match rate and fan-out recorded.
4. **Assess data quality**
   Six dimensions measured per column. Findings carry severity and affected row counts.
5. **Build dataset**
   Grain declared and enforced. Measure tables aggregated before joining.
6. **Transform data**
   Audited operations only, each showing rows before and after.
7. **Validate dataset** — Soft gate
   Eleven checks. Warnings flag and pass; failures stop the run.
8. **Generate documentation**
   Dictionary, dataset doc and lineage produced from the build record.
9. **Publish trusted dataset** — Hard gate · no override
   Source and target totals compared within tolerance. Outside tolerance, publication is refused.
10. **Prepare for analytics & AI**
    Readiness status issued: analytics, BI and model training, each gated on the checks above.

Stages 1–4 · Understand

Connection, discovery, relationship verification and quality assessment. Output is knowledge, not data. Nothing is written.

Stages 5–7 · Produce

Build, transform and validate. Every operation records rows before and after, and every AI proposal carries confidence, evidence, expected impact and risk.

Stages 8–10 · Prove

Documentation, reconciliation and publication. The reconciliation gate is the only step in the workflow with no override path.

SECTION 10

Key Platform Modules

Fifteen modules. Each entry states its purpose and the specific value it creates — not its feature list.

#	MODULE	PURPOSE	VALUE CREATED
01	**Data Source Connectivity**	Read-only, credential-vaulted connections to warehouses, databases, files and APIs.	One safe entry point. Credentials never reach a model; connections are least-privilege by default.
02	**Data Landscape Explorer**	A live map of domains, tables, columns and the relationships measured between them.	Replaces tribal knowledge with an accurate picture, including across different source systems.
03	**AI Discovery Center**	Schema, metadata and relationship discovery, with drafted business descriptions for review.	Compresses weeks of manual inventory into minutes, with a human confirming meaning.
04	**Data Health Center**	Six quality dimensions measured per column with a rules engine and severities.	Finds the problems nobody thought to test for, and quantifies quality for the board.
05	**Dataset Recommendation Engine**	Proposes candidate datasets with the tables required, estimated quality and confidence.	Starts work from a validated plan instead of a blank canvas.
06	**Visual Dataset Builder**	Assembles datasets with declared grain, verified joins and a live preview at every step.	Prevents the fan-out class of error at design time rather than discovering it in production.
07	**Transformation Studio**	A fixed library of audited operations with before and after shown for each.	No generated code reaches the data. Every change is reviewable and reversible.
08	**SQL Copilot**	Natural language to SQL composed only from verified join paths, with plan and impact shown before execution.	Speeds query writing while making the consequences visible before anything runs.
09	**Validation Center**	Completeness, consistency, freshness, integrity and an overall score, each explainable.	Converts "the data looks fine" into a figure that can be tracked and reported.
10	**Documentation Center**	Data dictionary, dataset documentation and glossary generated from the build record.	Documentation that cannot drift, produced at zero marginal engineer time.

Key Platform Modules

#	MODULE	PURPOSE	VALUE CREATED
11	**Data Lineage**	Column-level lineage from source through raw, clean and curated layers to datasets, reports and models.	Answers "where did this number come from" in seconds instead of days, with every step clickable.
12	**Monitoring**	Pipeline health, freshness, volume, schema change and quality drift, with alerts routed to owners.	Problems are found by the platform rather than by a business user disputing a figure.
13	**Impact Analysis**	Before a change ships: which datasets, pipelines, reports and downstream systems are affected.	Turns a change from a gamble into a planned action with a known blast radius.
14	**Agent Control Center**	Status, confidence and task history for every AI agent operating in the platform.	Makes AI activity observable and accountable rather than opaque.
15	**Governance Center**	Audit log, access control, policies, sensitive-data classification and compliance status.	Evidence for an auditor produced as a by-product of normal work.

Two modules that define the product

Reconciliation, inside Validation

Every published measure is compared against its source within a declared tolerance. A variance outside tolerance blocks publication and no reviewer can override it. Most tools report a variance; refusing to release on one is a different product behaviour, and it is the behaviour customers remember.

The Business Knowledge Layer

Company-specific definitions — what "active customer" means here, how default is measured, which risk bands apply — stored once, versioned, owned, and read by the platform *before* it reads the schema. This is what makes the AI sound like it understands the business rather than the column names.

Deliberate sequencing. Modules 1, 2, 4, 6, 9 and 10 constitute the first shippable product. The remainder are staged behind evidence of adoption. Building all fifteen before the first customer is how products in this category fail: eighteen things each seventy per cent complete.

SECTION 11

AI Agent Strategy

Seven narrow agents, each with a defined input, a defined output and a confidence score. None holds a database connection. They propose; a deterministic engine executes; a human approves.

AGENT	PURPOSE	INPUTS	OUTPUTS	BENEFIT
Discovery	Map schema, metadata and candidate relationships	Table and column names, types, null and distinct counts, declared constraints	Inventory, candidate relationships, drafted descriptions	Weeks of manual inventory reduced to minutes
Quality	Detect nulls, duplicates, outliers and type problems	Column profiles and distributions	Findings with severity, affected row counts and proposed fixes	Finds the problems nobody wrote a test for
Join	Propose joins and measure how well they hold	Declared keys, name similarity, measured value overlap	Verified paths with match rate, confidence and fan-out ratio	Only this agent may mark a path verified
Transformation	Decide which cleaning and shaping steps are needed	Quality findings, target grain, business definitions	An ordered plan drawn from a fixed library of audited operations	Cannot write arbitrary code — every step is reviewable
Documentation	Produce dictionary, descriptions and glossary entries	Catalogue, compiled plan, build record	Draft documentation for owner approval	Documentation that cannot drift from the pipeline
Monitoring	Watch pipelines, freshness, volume and drift	Run history, metric time series, schema snapshots	Alerts grouped into incidents with a probable cause	Correlates related signals instead of raising five separate alarms
Governance	Check policy and classify sensitive data	Column profiles, naming patterns, policy definitions	Classification, policy findings, blocked actions	Prevents an unmasked share before it happens

What every agent is denied

No database connection. No file-system access. No ability to execute. Agents receive schema metadata, business definitions and measurements the platform has taken. Row values, customer names, balances, account numbers and sample dataframes are never in the payload — a boundary enforced in code and verified by an automated test suite.

Why narrow agents, in a fixed order

Quality runs before Join, because a join measured on dirty keys reports a false match rate. Transformation runs after Join, because the aggregation level depends on the fan-out Join measured. The sequence is engineered, not negotiated between agents at runtime — which makes behaviour reproducible and failures diagnosable.

SECTION 12

Competitive Landscape

Assessment current to late 2026; these products move quickly and should be re-verified before any external use of this page.

VENDOR	STRENGTHS	WEAKNESSES	HOW WE DIFFER
Atlan	Best-in-class experience, search, Slack and Jira integration, fast adoption	Describes data; does not build or verify it. No arithmetic guarantee.	We produce and prove the dataset rather than catalogue it
Collibra	Deep governance workflow, stewardship, regulatory credibility, enterprise sales reach	Long implementations, heavy to operate, low day-to-day engineer adoption	Governance as a by-product of building, not a parallel programme
Alation	Strong search, query-log mining, established with BI analysts	Tells you which table is popular, not whether the result is correct	Popularity is not correctness; we measure correctness
DataHub	Open source, free, extensible, strong column-level lineage	Requires a platform team to run; no verification or build capability	Usable without a dedicated platform team; we verify as well as trace
Informatica	Connector breadth, scale, legacy ETL and mainframe understanding, 30-year relationships	Cost, complexity, slow to adopt, limited AI-native workflow	Faster to value; we verify output rather than only executing movement
Databricks	Everything, if the customer is on Databricks. Unity Catalog is bundled.	Single-platform. Weak across a mixed estate of SQL Server, Oracle and Snowflake.	Platform-neutral across the mixed estates most enterprises actually run

CAPABILITY COMPARISON

CAPABILITY	ATLAN	COLLIBRA	ALATION	DATAHUB	INFORMATICA	DATABRICKS	US
Catalogue & search	Strong	Strong	Strong	Good	Good	Good	Basic
Governance workflow	Good	Strong	Good	Basic	Good	Good	Planned
Connector breadth	Strong	Strong	Good	Good	Strong	Good	Narrow
Dataset building	No	No	No	No	Strong	Strong	Strong
Grain & fan-out safety	No	No	No	No	Manual	Manual	Core
Reconciliation as a blocking gate	No	No	No	No	Manual	No	Core
Decision provenance	Partial	Good	Partial	Partial	Partial	Partial	Core

Honest reading of this table. We are behind on four of seven rows, and the gaps — connectors, catalogue, governance workflow — are real and will be raised in every evaluation. Our case rests on two rows no competitor occupies. A deal is won by moving the conversation to those two rows, and lost by competing on the other five.

SECTION 13

Unique Value Proposition

> We are the only platform that will refuse to publish your dataset when the numbers do not add up.
>
> The single sentence a buyer should remember

Faster dataset creation

Discovery, profiling and join verification collapse from weeks to hours. The target is a 4–6× reduction in elapsed time for a typical multi-table dataset — achieved by removing waiting and rework, not by removing review.

AI-assisted engineering

AI narrows two thousand possible column pairings to twelve candidates worth a human's attention, drafts the query and drafts the documentation. The engineer's judgement is applied where it is scarce and valuable.

Data trust, evidenced

Every published dataset carries a reconciliation result, a quality score and a named approver. Trust stops being a feeling and becomes a record that survives an audit.

Documentation without effort

Generated from the build record at zero marginal engineer time, and structurally unable to drift from what the pipeline does.

Knowledge retention

Why a join exists, why a rule was added, why a suggestion was rejected — recorded at the moment of decision. When the engineer leaves, the reasoning stays. This is the quietest benefit and often the most valuable.

Impact known before change

What breaks if this column changes is answered before the change ships, with the affected datasets, reports and models named.

WHERE THE VALUE LANDS, BY ROLE

ROLE	BEFORE	AFTER
Data Engineer	Four to six weeks per dataset; reconciliation skipped under deadline	Days; reconciliation automatic and non-negotiable
Analytics Engineer	Tests catch only anticipated failures	Grain, fan-out and contract breaks caught structurally
Business Analyst	Cannot tell authoritative tables from abandoned copies	Certified datasets with owner, definition and quality attached
Chief Data Officer	Quality is an assertion	Quality, coverage and reconciliation are reportable numbers
Compliance	Evidence reconstructed after the fact	Who, what, when and why captured at the moment of action

SECTION 14

Success Metrics

Targets below are the hypotheses this product will be judged against. They are written to be measurable in a customer's own environment, and to be falsifiable.

PRODUCT OUTCOME METRICS

METRIC	BASELINE TODAY	12-MONTH TARGET	HOW IT IS MEASURED
Dataset creation time	4–6 weeks	5–7 days	Elapsed days from request raised to dataset published
Documentation coverage	< 15% of columns	100% of published	Published columns carrying an owner-approved definition
Reconciliation coverage	Ad hoc, often skipped	100%	Published datasets with a passing reconciliation on record
Data quality score	Not measured	+15 points	Composite across six dimensions, tracked monthly per domain
Dataset reuse rate	Near zero	≥ 40%	New datasets cloned from an existing one rather than built fresh
Engineer time recovered	—	20–30%	Self-reported time on discovery, verification and documentation
Silent data errors	Discovered by business users	Caught pre-publish	Count of incidents raised by consumers rather than by the platform

ADOPTION AND COMMERCIAL METRICS

WEEKLY ACTIVE ≥ 60% of licensed data engineers, by month 6

PILOT CONVERSION ≥ 40% of paid pilots converting to annual contract

TIME TO FIRST VALUE < 5 days connection to first reconciled dataset

NET RETENTION ≥ 110% by year two, via domain expansion

A metric we will not claim. Headcount reduction. The platform recovers engineer time and redirects it; it does not remove roles. Promising headcount savings produces a business case that fails at renewal review, and the credibility cost is not recoverable.

SECTION 15

Product Roadmap Vision

Five phases. Each is gated on evidence from the one before it, not on a calendar. Phases 4 and 5 are strategic direction rather than commitments.

Phase 1 — Trusted Dataset Platform

Foundation · ship first

Connect, discover, profile, assess quality, build with correct grain, reconcile as a blocking gate, generate documentation and lineage. Plus the unglamorous prerequisites: authentication, users and roles. This is a complete, sellable product on its own — and the only phase that must be right before anything else matters.

Phase 2 — Governance & Knowledge Layer

Unlocks regulated buyers

Business Knowledge Layer, audit logging, data contracts and schema monitoring, full source-to-model lineage, and the Agent Control Center. This is the phase that moves the product from a useful tool to one a compliance function will approve.

Phase 3 — Enterprise Intelligence Layer

After real usage exists

Monitoring and alerting, impact analysis, root cause investigation, dataset marketplace and enterprise catalogue, broader connector coverage. Every item here needs real operational history to be worth anything — building it against assumptions wastes the investment.

Phase 4 — Advanced AI Agents

Directional

Agents that learn an organisation's conventions over time, propose remediation with predicted impact, and reason across the accumulated decision record. Credible only once the decision record has real depth.

Phase 5 — Autonomous Data Operations

Long horizon

Self-healing pipelines, automatic contract renegotiation, continuous dataset certification. We state plainly that human approval remains in the loop for anything that changes a published figure. Autonomy here means less routine work, not an absence of accountability.

The discipline this roadmap encodes. Phase 1 is narrow on purpose. A platform that does six things excellently will beat one that does thirty adequately in every enterprise evaluation — and the second is how small teams exhaust their runway.

SECTION 16

The intelligence layer for enterprise data

Enterprises have solved storage. They have solved compute. They have not solved understanding — and understanding is what every analytics and AI initiative silently depends on.

Data Intelligence OS is the layer that knows what the data means, proves whether it is right, and remembers why every decision about it was made. It does not replace the warehouse, the pipeline or the catalogue. It makes their output trustworthy.

Trusted data

Not asserted. Reconciled, evidenced and signed.

Faster decisions

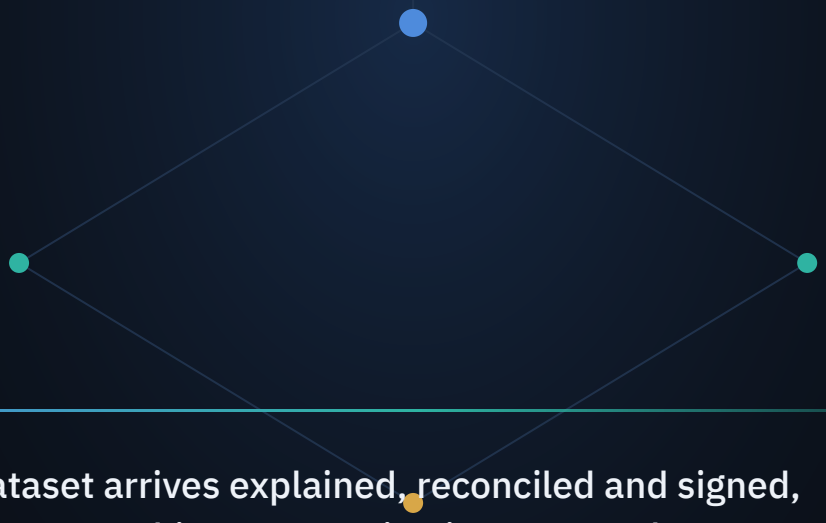
Weeks of discovery removed from every question the business asks.

AI-assisted operations

Judgement applied where it is scarce; the routine work automated.

Knowledge preserved

The reasoning outlives the person who built the dataset.



When every dataset arrives explained, reconciled and signed, data stops being something an organisation argues about — and becomes something it decides with.

DOCUMENT	DATE	PREPARED BY	CLASSIFICATION
Product Vision v1.0	October 2026	PyrunAI Services LLP	Confidential